

POSIX Order Write Test Report

Purpose

The report explains the requirements and design of the POSIX Order Write test in the SWMR project. It then shows the result of the tests in different systems.

Requirements

The write order test should verify that the write order is strictly consistent. The SWMR feature requires that the order of write is strictly consistent. The SWMR updates to the data structures in the file are essentially implementing a "lock-free" or "wait-free" algorithm in updating the data structure on disk. See: https://en.wikipedia.org/wiki/Non-blocking_synchronization, for example. In those algorithms, the order of updates to the data structure is critical

The other reader processes will try to read the address value at the offset 0. The value is initially NULL(0). When it changes to non-zero, it signifies the writer process has written all the chain-link blocks and they are ready for the reader processes to access.

If the system, in which the writer and reader processes run, adhere to the order of write, the readers will always get all chain-linked blocks correctly. If the order of write is not maintained, some reader processes may find unexpected block data.

Results

AIX hosts with GPFS

The machines run AIX 5.3 OS and are as a box of “blades”. All blades share the access to the GPFS filesystem.

The test, including both writer and reader, running in the same or separated “blades”, passed all runs up to 1,000,000 linked blocks, resulting in datafiles ~2GB big.

A side note: when in separated blades, the write time is 7.6 seconds but the reader time is 69 seconds. It is speculated that the GPFS system may be doing some “kernel” buffering to make write time to appear smaller.

Linux hosts with GPFS

The above mentioned AIX system site also has Linux hosts that share the same GPFS system. The Linux hosts are 64bits system.

The POSIX write order test is run in separated Linux machines using the same GPFS file system. The test, including both write and reader, in the same or separated hosts, passed all tests up to 1,000,000 linked blocks which resulted in approximately 2GB size file.

A side note: the writer took 8.5 sec to write the 2GB file but the reader took 83 seconds to read them. Write speed is 10 times faster than read speed--Kernel memory is in play here. Nevertheless, IBM GPFS adheres to the write order correctly.

Linux Cluster with Lustre file system

The POSIX write order test is run in a remote Linux Cluster with a Lustre file system. The test passed with small size files such as 200MB in size. But when larger number of linked blocks (e.g. -n 500000 => 500,000 linked blocks resulting in file size about 900MB) and the writer and reader were in separated machines, the reader detected errors in data it read back. The exact cause of the failure is not known yet but it fails for bigger file sizes over separated machines.

Summary

This is a summary of the test results in different system using different the GPFS or Lustre file system.

System	Same host/machine	Separated hosts/machines
AIX with GPFS	Passed in all file sizes up to 2GB.	Passed in all file sizes up to 2GB.
Linux with GPFS	Passed in all file sizes up to 2GB.	Passed in all file sizes up to 2GB.
Linux with Lustre	N/A	Passed in small file sizes but failed with large files (e.g., 900MB)